

## 6.824 2017 Lecture 1: Introduction

### 6.824: Distributed Systems Engineering

What is a distributed system?

- multiple cooperating computers
- big databases, P2P file sharing, MapReduce, DNS, &c
- lots of critical infrastructure is distributed!

Why distributed?

- to connect physically separate entities
- to achieve security via isolation
- to tolerate faults via replication
- to scale up throughput via parallel CPUs/mem/disk/net

But:

- complex: many concurrent parts
- must cope with partial failure
- tricky to realize performance potential

Why take this course?

- interesting -- hard problems, powerful solutions
- used by real systems -- driven by the rise of big Web sites
- active research area -- lots of progress + big unsolved problems
- hands-on -- you'll build serious systems in the labs

### COURSE STRUCTURE

<http://pdos.csail.mit.edu/6.824>

Course staff:

- Robert Morris, lecturer
- Frans Kaashoek, lecturer
- Lara Araujo, TA
- Anish Athalye, TA
- Srivatsa Bhat, TA
- Daniel Ziegler, TA

Course components:

- lectures
- readings
- two exams
- labs
- final project (optional)

Lectures:

big ideas, paper discussion, and labs

Readings:

research papers, some classic, some new  
the papers illustrate key ideas and important details  
many lectures focus on the papers  
please read papers before class!  
each paper has a short question for you to answer  
and you must send us a question you have about the paper  
submit question&answer by 10pm the night before

Mid-term exam in class; final exam during finals week

Lab goals:

deeper understanding of some important techniques  
experience with distributed programming  
first lab is due a week from Friday  
one per week after that for a while

Lab 1: MapReduce

Lab 2: replication for fault-tolerance using Raft

Lab 3: fault-tolerant key/value store

Lab 4: sharded key/value store

Optional final project at the end, in groups of 2 or 3.

The final project substitutes for Lab 4.

You think of a project and clear it with us.

Lab grades depend on how many test cases you pass

we give you the tests, so you know whether you'll do well  
careful: if it often passes, but sometimes fails,  
chances are it will fail when we run it

Debugging the labs can be time-consuming

start early  
come to TA office hours  
ask questions on Piazza

MAIN TOPICS

This is a course about infrastructure, to be used by applications.

About abstractions that hide distribution from applications.

Three big kinds of abstraction:

Storage.  
Communication.  
Computation.

[diagram: users, application servers, storage servers]

A couple of topics come up repeatedly.

Topic: implementation

RPC, threads, concurrency control.

Topic: performance

The dream: scalable throughput.

$N$ x servers  $\rightarrow$   $N$ x total throughput via parallel CPU, disk, net.

So handling more load only requires buying more computers.

Scaling gets harder as  $N$  grows:

Load im-balance, stragglers.

Non-parallelizable code: initialization, interaction.

Bottlenecks from shared resources, e.g. network.

Topic: fault tolerance

1000s of servers, complex net  $\rightarrow$  always something broken

We'd like to hide these failures from the application.

We often want:

Availability -- app can keep using its data despite failures

Durability -- app's data will come back to life when failures are repaired

Big idea: replicated servers.

If one server crashes, client can proceed using the other(s).

Topic: consistency

General-purpose infrastructure needs well-defined behavior.

E.g. "Get( $k$ ) yields the value from the most recent Put( $k,v$ )."

Achieving good behavior is hard!

"Replica" servers are hard to keep identical.

Clients may crash midway through multi-step update.

Servers crash at awkward moments, e.g. after executing but before replying.

Network may make live servers look dead; risk of "split brain".

Consistency and performance are enemies.

Consistency requires communication, e.g. to get latest Put().

"Strong consistency" often leads to slow systems.

High performance often imposes "weak consistency" on applications.

People have pursued many design points in this spectrum.

CASE STUDY: MapReduce

Let's talk about MapReduce (MR) as a case study

MR is a good illustration of 6.824's main topics

and is the focus of Lab 1

## MapReduce overview

context: multi-hour computations on multi-terabyte data-sets  
e.g. analysis of graph structure of crawled web pages  
only practical with 1000s of computers  
often not developed by distributed systems experts  
distribution can be very painful, e.g. coping with failure  
overall goal: non-specialist programmers can easily split  
data processing over many servers with reasonable efficiency.  
programmer defines Map and Reduce functions  
sequential code; often fairly simple  
MR runs the functions on 1000s of machines with huge inputs  
and hides details of distribution

## Abstract view of MapReduce

input is divided into M files  
Input1 -> Map -> a,1 b,1 c,1  
Input2 -> Map ->                   b,1  
Input3 -> Map -> a,1                   c,1  
                                          |            |            |  
                                                          |                    -> Reduce -> c,2  
                                                          -----> Reduce -> b,2

MR calls Map() for each input file, produces set of k2,v2  
"intermediate" data  
each Map() call is a "task"  
MR gathers all intermediate v2's for a given k2,  
and passes them to a Reduce call  
final output is set of <k2,v3> pairs from Reduce()  
stored in R output files

## Example: word count

input is thousands of text files  
Map(k, v)  
  split v into words  
  for each word w  
    emit(w, "1")  
Reduce(k, v)  
  emit(len(v))

## MapReduce hides many painful details:

starting s/w on servers  
tracking which tasks are done  
data movement  
recovering from failures

## MapReduce scales well:

N computers gets you Nx throughput.

Assuming M and R are  $\geq N$  (i.e. lots of input files and output keys).

Maps()s can run in parallel, since they don't interact.

Same for Reduce()s.

So you can get more throughput by buying more computers.

Rather than special-purpose efficient parallelizations of each application.

Computers are cheaper than programmers!

What will likely limit the performance?

We care since that's the thing to optimize.

CPU? memory? disk? network?

In 2004 authors were limited by "network cross-section bandwidth".

[diagram: servers, tree of network switches]

Note all data goes over network, during Map->Reduce shuffle.

Paper's root switch: 100 to 200 gigabits/second

1800 machines, so 55 megabits/second/machine.

Small, e.g. much less than disk or RAM speed.

So they cared about minimizing movement of data over the network.

(Datacenter networks are much faster today.)

More details (paper's Figure 1):

master: gives tasks to workers; remembers where intermediate output is

M Map tasks, R Reduce tasks

input stored in GFS, 3 copies of each Map input file

all computers run both GFS and MR workers

many more input tasks than workers

master gives a Map task to each worker

hands out new tasks as old ones finish

Map worker hashes intermediate keys into R partitions, on local disk

no Reduce calls until all Maps are finished

master tells Reducers to fetch intermediate data partitions from Map workers

Reduce workers write final output to GFS (one file per Reduce task)

How does detailed design reduce effect of slow network?

Map input is read from GFS replica on local disk, not over network.

Intermediate data goes over network just once.

Map worker writes to local disk, not GFS.

Intermediate data partitioned into files holding many keys.

Big network transfers are more efficient.

How do they get good load balance?

Critical to scaling -- bad for N-1 servers to wait for 1 to finish.

But some tasks likely take longer than others.

Solution: many more tasks than workers.

Master hands out new tasks to workers who finish previous tasks.  
So no task is so big it dominates completion time (hopefully).  
So faster servers do more work than slower ones, finish abt the same  
time.

What about fault tolerance?

I.e. what if a server crashes during a MR job?

Hiding failures is a huge part of ease of programming!

Why not re-start the whole job from the beginning?

MR re-runs just the failed Map()s and Reduce()s.

MR requires them to be pure functions:

they don't keep state across calls,

they don't read or write files other than expected MR  
inputs/outputs,

there's no hidden communication among tasks.

So re-execution yields the same output.

The requirement for pure functions is a major limitation of

MR compared to other parallel programming schemes.

But it's critical to MR's simplicity.

Details of worker crash recovery:

\* Map worker crashes:

master sees worker no longer responds to pings

crashed worker's intermediate Map output is lost

but is likely needed by every Reduce task!

master re-runs, spreads tasks over other GFS replicas of input.

some Reduce workers may already have read failed worker's intermediate  
data.

here we depend on functional and deterministic Map()!

master need not re-run Map if Reduces have fetched all intermediate  
data

though then a Reduce crash would then force re-execution of failed  
Map

\* Reduce worker crashes.

finished tasks are OK -- stored in GFS, with replicas.

master re-starts worker's unfinished tasks on other workers.

\* Reduce worker crashes in the middle of writing its output.

GFS has atomic rename that prevents output from being visible until  
complete.

so it's safe for the master to re-run the Reduce tasks somewhere else.

Other failures/problems:

\* What if the master gives two workers the same Map() task?

perhaps the master incorrectly thinks one worker died.

it will tell Reduce workers about only one of them.

\* What if the master gives two workers the same Reduce() task?

- they will both try to write the same output file on GFS!  
atomic GFS rename prevents mixing; one complete file will be visible.
- \* What if a single worker is very slow -- a "straggler"?  
perhaps due to flakey hardware.  
master starts a second copy of last few tasks.
  - \* What if a worker computes incorrect output, due to broken h/w or s/w?  
too bad! MR assumes "fail-stop" CPUs and software.
  - \* What if the master crashes?

For what applications \*doesn't\* MapReduce work well?

Not everything fits the map/shuffle/reduce pattern.

Small data, since overheads are high. E.g. not web site back-end.

Small updates to big data, e.g. add a few documents to a big index

Unpredictable reads (neither Map nor Reduce can choose input)

Multiple shuffles, e.g. page-rank (can use multiple MR but not very efficient)

More flexible systems allow these, but more complex model.

## Conclusion

MapReduce single-handedly made big cluster computation popular.

- Not the most efficient or flexible.

+ Scales well.

+ Easy to program -- failures and data movement are hidden.

These were good trade-offs in practice.

We'll see some more advanced successors later in the course.

Have fun with the lab!